

Data Loading and Cleaning

first i set option to display all columns dropping duplicated rows

replacing empty value with numpy nan

checking for null ways in any column

dropping rows which contain null values

deleting one column which is unnecessary $\text{loan_percent_income} = (\text{loan_amnt} / \text{person_income}) * 100$

Outliers

beyond age=84 there are ages like 144 and 123 which are clearly outliers

there are some cases which have retirement time greater than 100

for less memory usage and faster training i converted dtype from float64 to float32 and int64 to int32

separating output labels from features

Encoding by Pipeline

Using pipeline

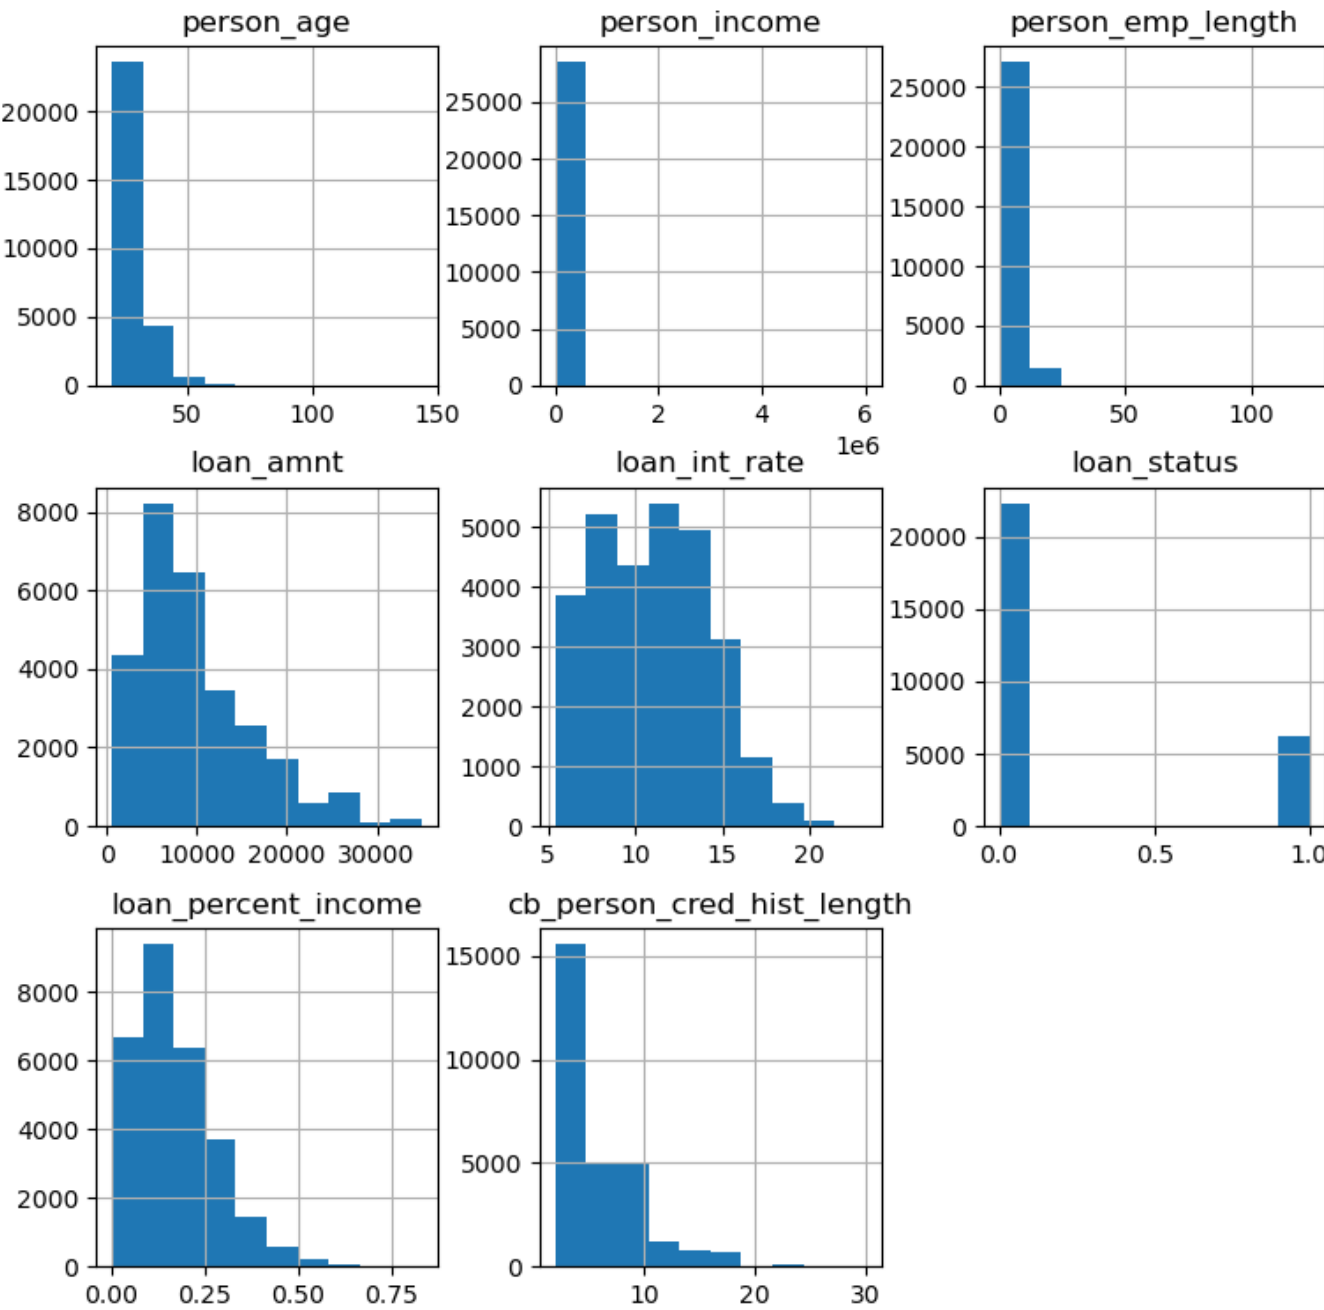
If you run the pipeline on the full dataset before splitting, steps like scaling or encoding will see the test data.

This causes data leakage, leading to unrealistic, overly optimistic model performance.

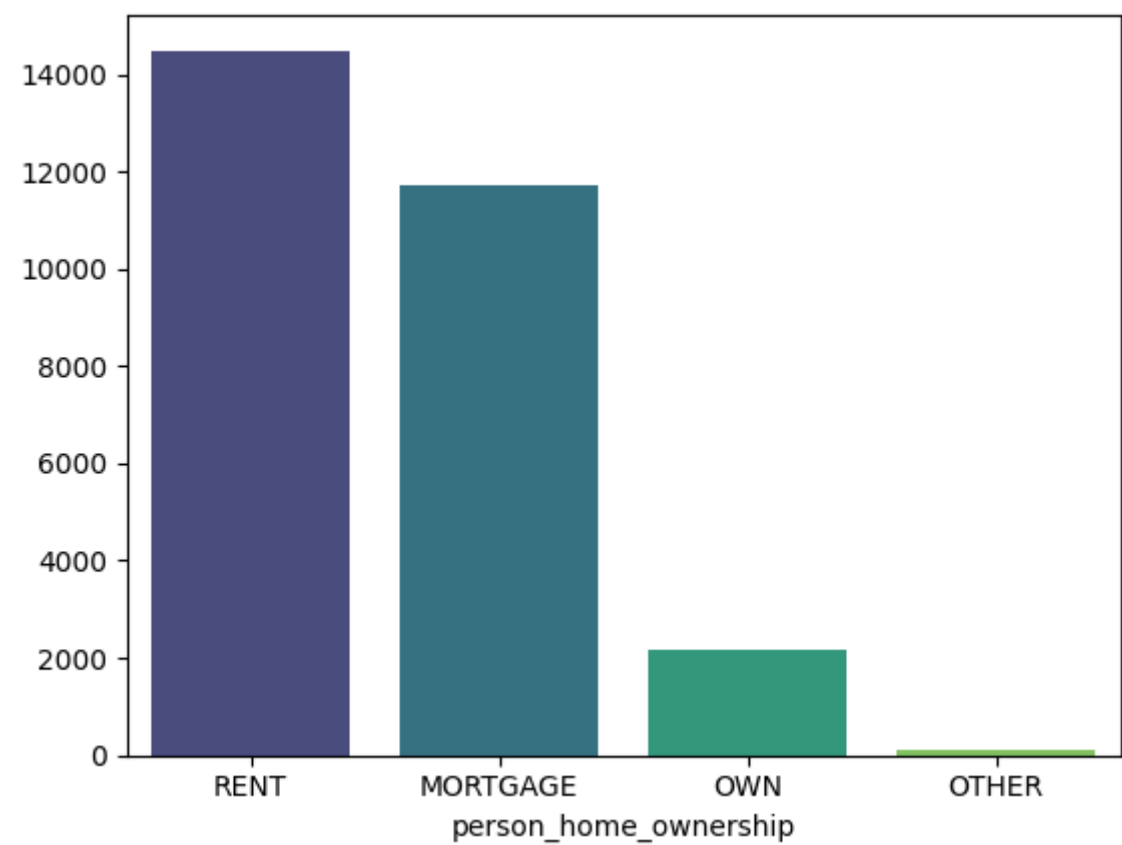
Saving preprocessed dataset

Visualizations

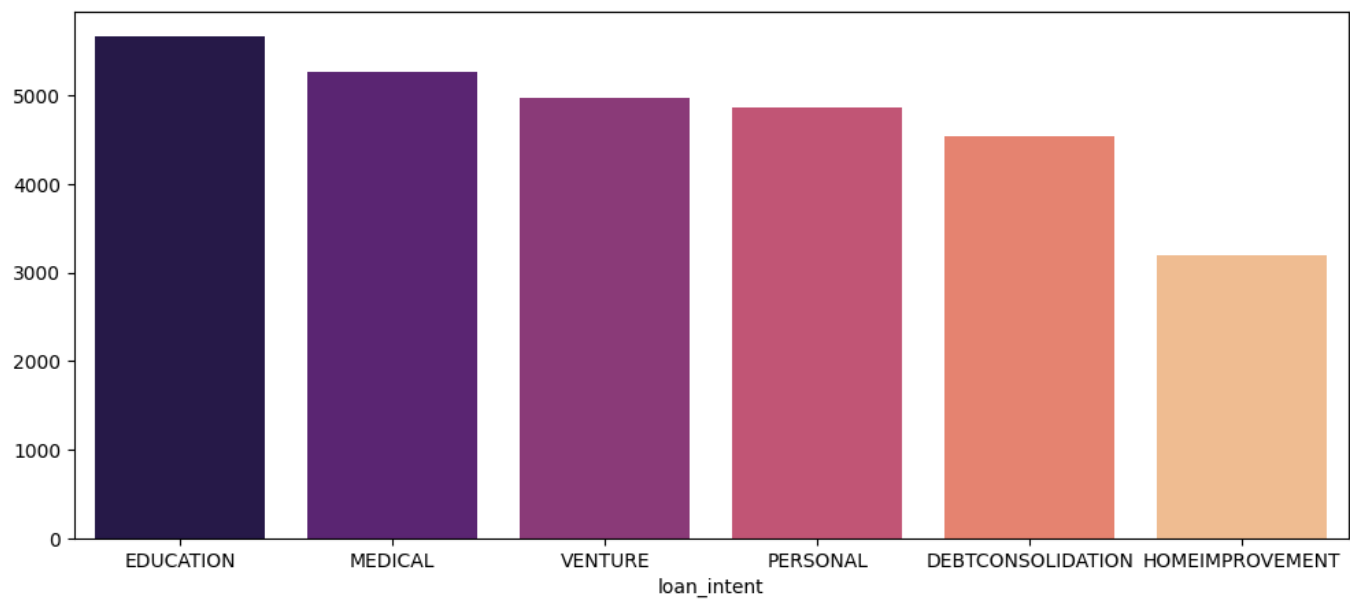
Hist plot



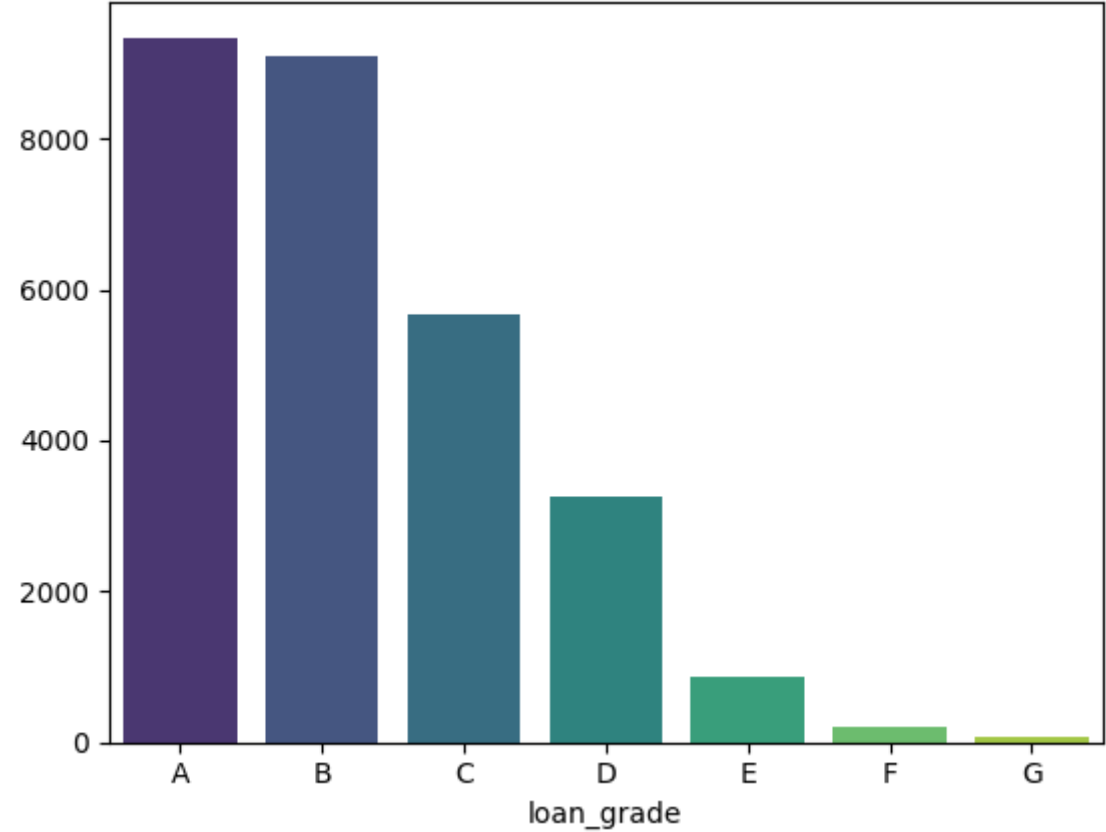
person_home_ownership



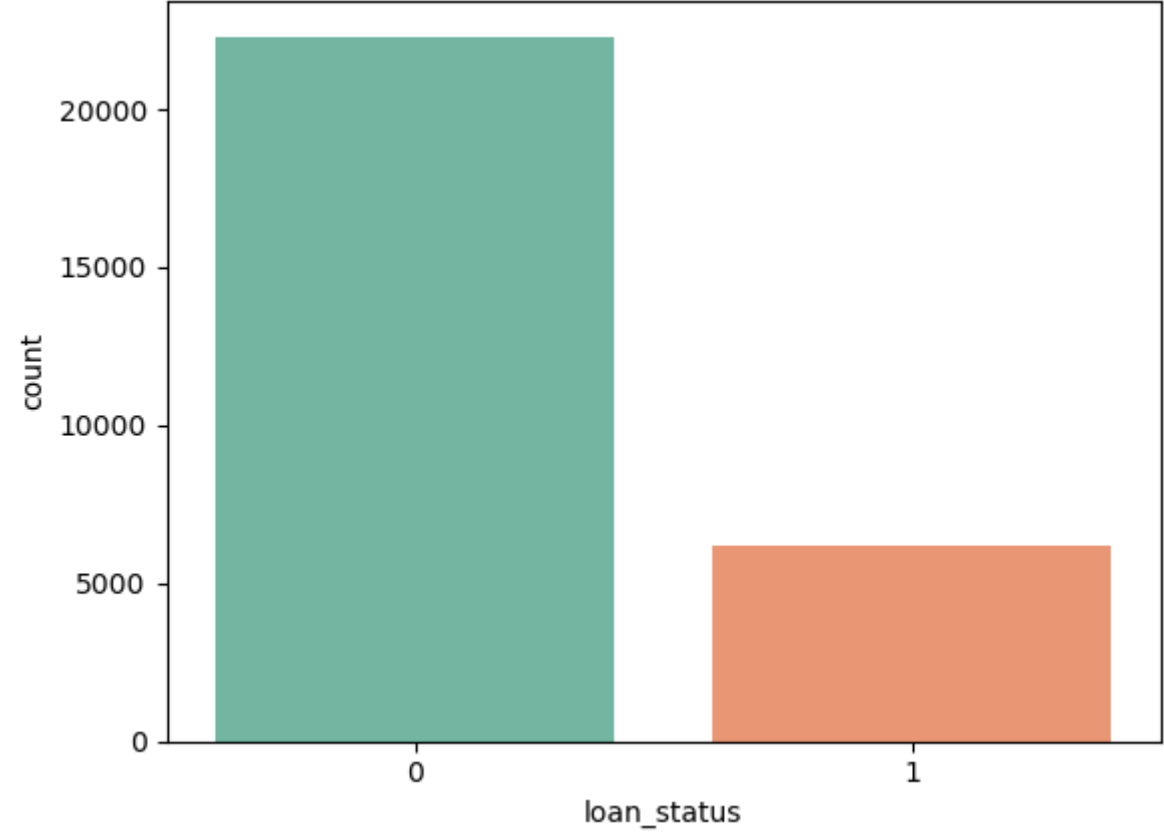
loan_intent



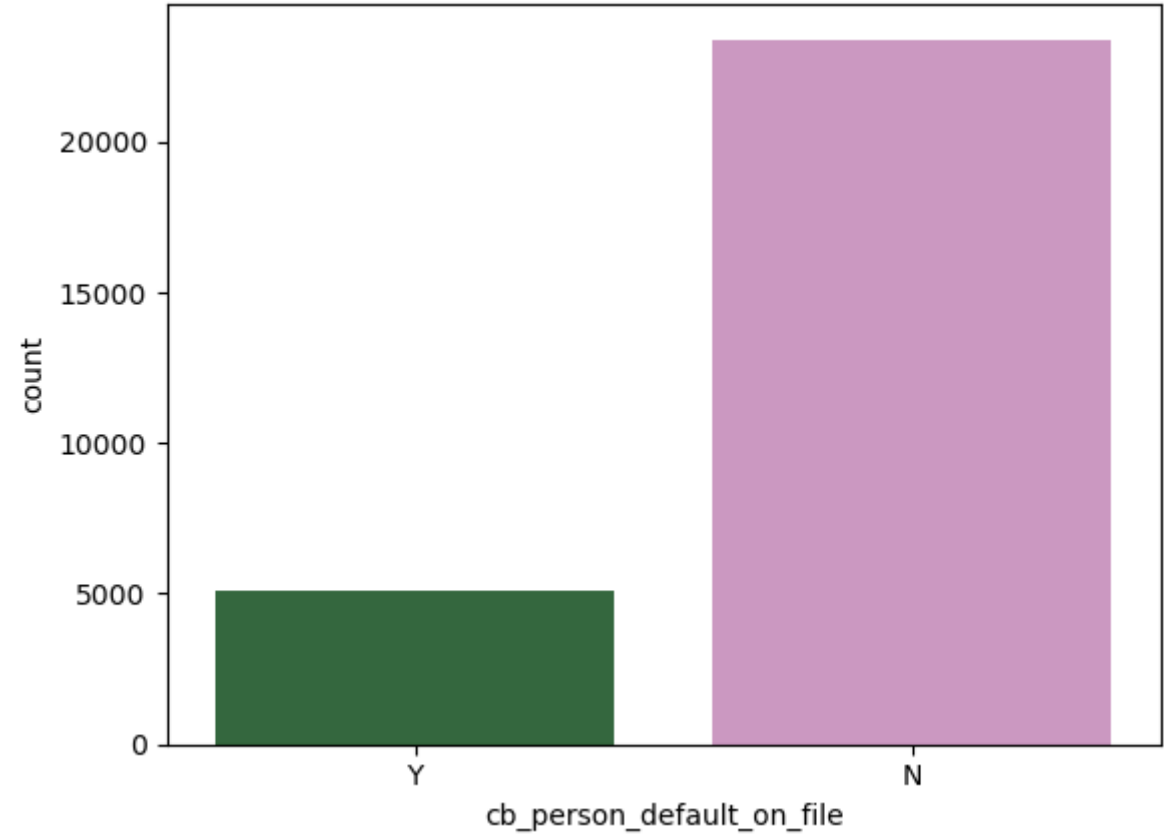
loan_grade

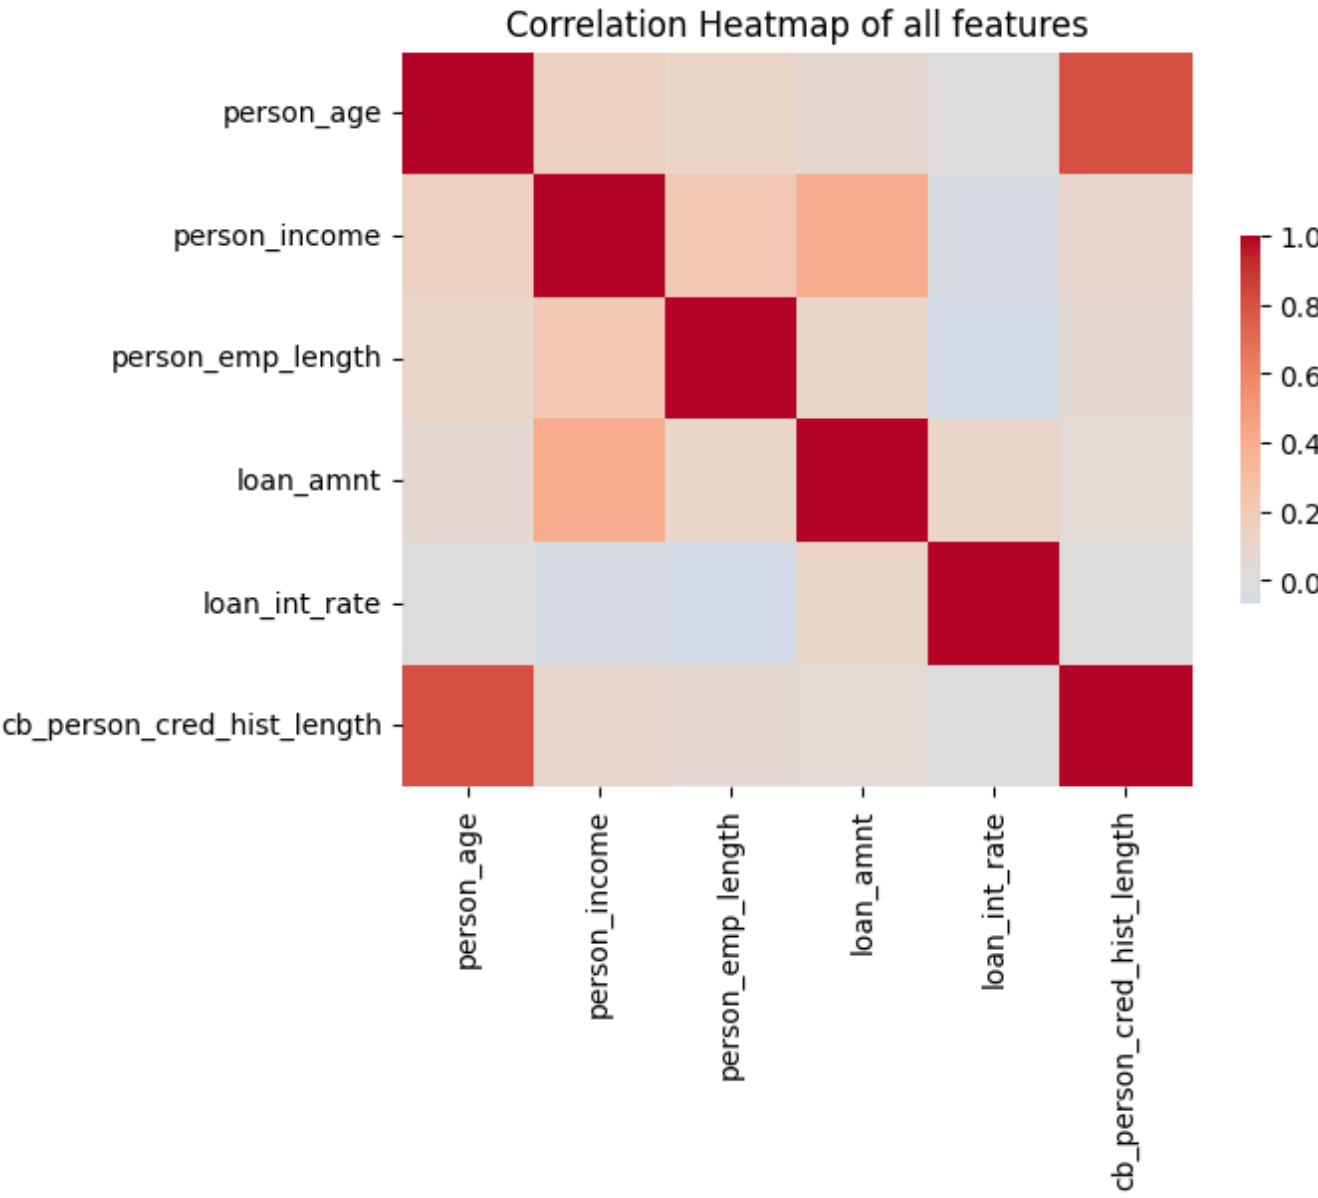


loan_status



cb_person_default_on_file





Logistic Regression

Here, mostly all are having best_f1=0.61 so default is good

```
solver=lbfgs, penalty=l2 and C=0.01
      precision    recall  f1-score   support
0         0.92        0.76        0.84       4462
1         0.48        0.77        0.59       1237
best_f1: 0.60
best_thresh: 0.68

solver=liblinear, penalty=l2 and C=0.01
      precision    recall  f1-score   support
0         0.92        0.76        0.83       4462
1         0.47        0.78        0.59       1237
best_f1: 0.60
best_thresh: 0.64
```

```
solver=liblinear, penalty=l1 and C=0.01
      precision    recall  f1-score   support
0         0.92      0.76      0.83     4462
1         0.47      0.78      0.59     1237
best_f1: 0.61
best_thresh: 0.65
```

```
solver=saga, penalty=l2 and C=0.01
      precision    recall  f1-score   support
0         0.92      0.76      0.84     4462
1         0.48      0.77      0.59     1237
best_f1: 0.60
best_thresh: 0.68
```

```
solver=saga, penalty=l1 and C=0.01
      precision    recall  f1-score   support
0         0.92      0.76      0.84     4462
1         0.48      0.77      0.59     1237
best_f1: 0.61
best_thresh: 0.65
```

```
solver=lbfgs, penalty=l2 and C=0.1
      precision    recall  f1-score   support
0         0.92      0.77      0.84     4462
1         0.48      0.77      0.59     1237
best_f1: 0.61
best_thresh: 0.63
```

```
solver=liblinear, penalty=l2 and C=0.1
      precision    recall  f1-score   support
0         0.92      0.76      0.84     4462
1         0.47      0.77      0.59     1237
best_f1: 0.61
best_thresh: 0.63
```

```
solver=liblinear, penalty=l1 and C=0.1
      precision    recall  f1-score   support
0         0.92      0.76      0.84     4462
1         0.47      0.77      0.59     1237
best_f1: 0.61
best_thresh: 0.63
```

```
solver=saga, penalty=l2 and C=0.1
      precision    recall  f1-score   support
0         0.92      0.76      0.84     4462
1         0.47      0.77      0.59     1237
```

```
best_f1: 0.61
best_thresh: 0.67

solver=saga, penalty=l1 and C=0.1
      precision    recall  f1-score   support
0         0.92        0.77        0.84        4462
1         0.48        0.77        0.59        1237
best_f1: 0.61
best_thresh: 0.67

solver=lbfgs, penalty=l2 and C=1
      precision    recall  f1-score   support
0         0.92        0.76        0.84        4462
1         0.48        0.77        0.59        1237
best_f1: 0.61
best_thresh: 0.63

solver=liblinear, penalty=l2 and C=1
      precision    recall  f1-score   support
0         0.92        0.76        0.84        4462
1         0.47        0.77        0.59        1237
best_f1: 0.61
best_thresh: 0.63

solver=liblinear, penalty=l1 and C=1
      precision    recall  f1-score   support
0         0.92        0.76        0.84        4462
1         0.48        0.77        0.59        1237
best_f1: 0.61
best_thresh: 0.63

solver=saga, penalty=l2 and C=1
      precision    recall  f1-score   support
0         0.92        0.76        0.84        4462
1         0.47        0.77        0.59        1237
best_f1: 0.61
best_thresh: 0.63

solver=saga, penalty=l1 and C=1
      precision    recall  f1-score   support
0         0.92        0.76        0.84        4462
1         0.47        0.77        0.59        1237
best_f1: 0.61
best_thresh: 0.63
```

Decision Tree

max_depth=10 seems to be optimal for decision tree

```
max_depth= None, best_f1: 0.73, best_thresh: 0.01
max_depth= 3, best_f1: 0.58, best_thresh: 0.48
max_depth= 5, best_f1: 0.72, best_thresh: 0.58
max_depth= 7, best_f1: 0.79, best_thresh: 0.81
max_depth= 10, best_f1: 0.80, best_thresh: 0.81
```

```
max_depth= 10, best_f1: 0.80, best_thresh: 0.81
max_depth= 15, best_f1: 0.78, best_thresh: 0.96
max_depth= 20, best_f1: 0.75, best_thresh: 0.92
max_depth= 40, best_f1: 0.73, best_thresh: 0.01
```

RandomForestClassifier

```
for estimators in [100,200]:
    for max_depth in [7,10,15,20]:
        for max_features in ["sqrt","log2"]:
```

```
max_depth=7, estimators=100 and max_features=sqrt
best_f1: 0.74, best_thresh: 0.60
```

```
max_depth=7, estimators=100 and max_features=log2
best_f1: 0.74, best_thresh: 0.60
```

```
max_depth=10, estimators=100 and max_features=sqrt
best_f1: 0.79, best_thresh: 0.62
```

```
max_depth=10, estimators=100 and max_features=log2
best_f1: 0.79, best_thresh: 0.62
```

```
max_depth=15, estimators=100 and max_features=sqrt
best_f1: 0.81, best_thresh: 0.55
```

```
max_depth=15, estimators=100 and max_features=log2
best_f1: 0.81, best_thresh: 0.55
```

```
max_depth=20, estimators=100 and max_features=sqrt
best_f1: 0.81, best_thresh: 0.50
```

```
max_depth=20, estimators=100 and max_features=log2
best_f1: 0.81, best_thresh: 0.50
```

```
max_depth=7, estimators=200 and max_features=sqrt
best_f1: 0.74, best_thresh: 0.59
```

```
max_depth=7, estimators=200 and max_features=log2
best_f1: 0.74, best_thresh: 0.59
```

```
max_depth=10, estimators=200 and max_features=sqrt
best_f1: 0.79, best_thresh: 0.62
```

```
max_depth=10, estimators=200 and max_features=log2
best_f1: 0.79, best_thresh: 0.62
```

```
max_depth=15, estimators=200 and max_features=sqrt
best_f1: 0.81, best_thresh: 0.53
```

```
max_depth=15, estimators=200 and max_features=log2
best_f1: 0.81, best_thresh: 0.53
```

```
max_depth=20, estimators=200 and max_features=sqrt
best_f1: 0.81, best_thresh: 0.48
```

```
max_depth=20, estimators=200 and max_features=log2
best_f1: 0.81, best_thresh: 0.48
```

Here $n_estimators=100/200$ and $max_features=sqrt/log2$ have same performance thus we will fix these now

performance increases as $0.74-0.79-0.81-0.81$ for $max_depth= [7,10,15,20]$ thus we will search now in range (10-15)

(keeping fix $n_estimators=100$, $max_features=sqrt$)

```
for max_depth in [12,13,14,15]:
    for min_samples_split in [10,20,40]:
        for min_samples_leaf in [3,5,10]:
```

after training on this range best_f1 after threshold optimization was 0.80 for most of them

For these combinations although best_f1 was 0.81

```
max_depth=13, min_samples_split=10 and min_samples_leaf=5
```

```
max_depth=14, min_samples_split=10 and min_samples_leaf=5
```

```
max_depth=14, min_samples_split=10 and min_samples_leaf=3
```

```
max_depth=14, min_samples_split=20 and min_samples_leaf=3
```

```
max_depth=15, min_samples_split=10 and min_samples_leaf=3
```

```
max_depth=15, min_samples_split=10 and min_samples_leaf=5
```

```
max_depth=15, min_samples_split=20 and min_samples_leaf=3
```

```
max_depth=15, min_samples_split=20 and min_samples_leaf=5
```

Choosing this as our best model

```
RandomForestClassifier(max_depth=15,min_samples_leaf=5,min_samples_split=20,n_estimators=100,max_features="sqrt",random_state=42)
```

	precision	recall	f1-score	support
0	0.93	0.96	0.95	4462
1	0.83	0.75	0.79	1237

```
Confusion Matrix: [[4277 185]
                   [ 307 930]]
```

ROC AUC Score: 0.93

best_f1: 0.81

best_thresh: 0.57

Feature Selection

Now extracting feature importances from our best RandomForest model

Here, 14 features conserve 99% of feature importances
while 11 features out of 16 conserve 95% of feature importances

creating new training and test set based on top 14 and 11 features

Now training on both 11 and 14 features

With top 14 features

	precision	recall	f1-score	support
0	0.94	0.95	0.94	4462
1	0.82	0.76	0.79	1237

```
Confusion Matrix: [[4254 208]
                   [ 295 942]]
```

ROC AUC Score: 0.93

best_f1: 0.80

best_thresh: 0.57

With top 11 features

	precision	recall	f1-score	support
0	0.94	0.94	0.94	4462
1	0.78	0.77	0.78	1237

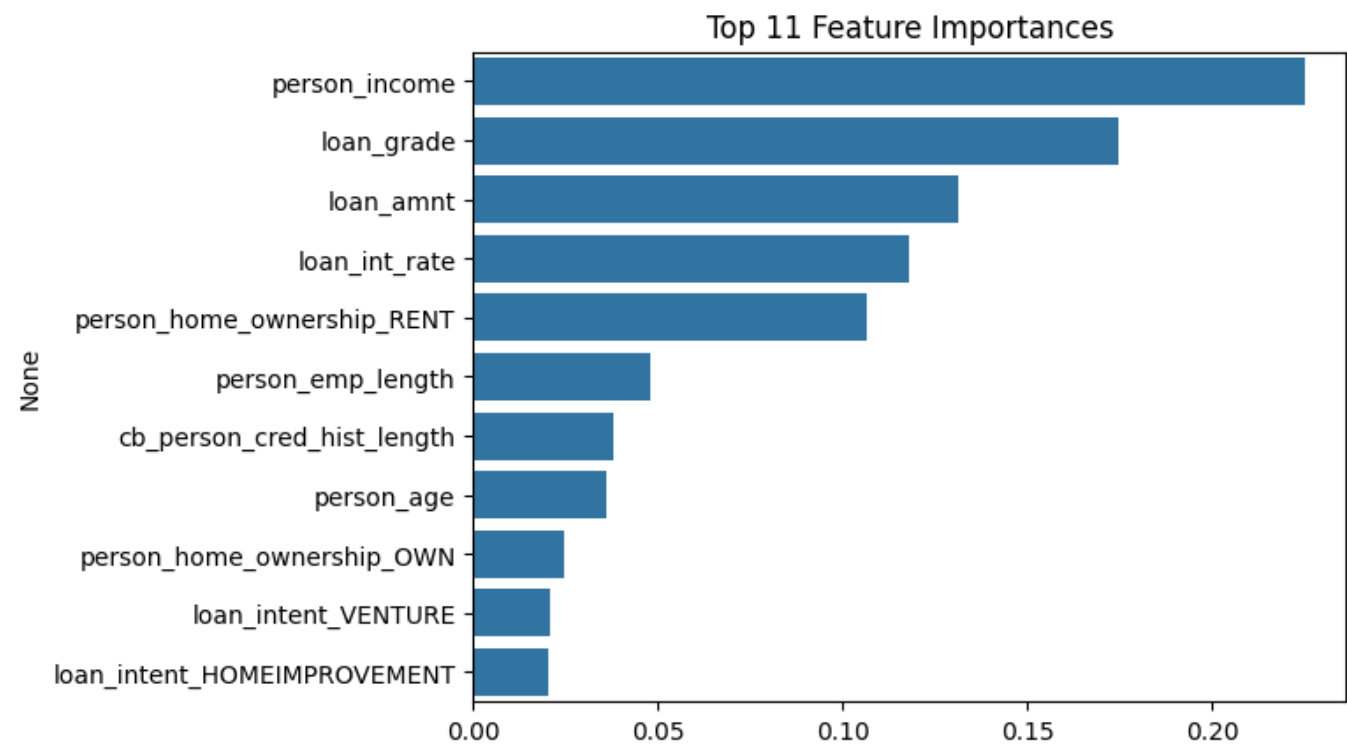
```
Confusion Matrix: [[4194 268]
                   [ 281 956]]
```

ROC AUC Score: 0.92

best_f1: 0.78

best_thresh: 0.50

Based on this results we can choose to pick 14 features rather than 11 features although, all 16 features had greater performance



XGBoostClassifier

```
n_estimators=500,  
learning_rate=0.05,  
max_depth=15,  
eval_metric="aucpr"
```

	precision	recall	f1-score	support
0	0.94	0.98	0.96	4462
1	0.92	0.76	0.83	1237

Confusion Matrix: $\begin{bmatrix} 4377 & 85 \\ 302 & 935 \end{bmatrix}$

ROC AUC Score: 0.94

best_f1: 0.83

best_thresh: 0.51

Using Early Stopping

Creating validation set from training set with ratio 0.2

Using same parameters with early_stopping_rounds=30

```
best_iteration= 329, best_score= 0.9889183509672792
      precision    recall  f1-score   support
0         0.93      0.98      0.96      4462
1         0.90      0.75      0.82      1237
```

```
Confusion Matrix: [[4363   99]
                   [ 305  932]]
```

```
ROC AUC Score: 0.94
```

```
best_f1: 0.83
```

```
best_thresh: 0.58
```

GridSearchCV

We can not use Early Stopping with GridSearchCV

```
"n_estimators":[500],
"max_depth":[15,20],
"gamma":[0,0.1,0.3],
"learning_rate":[0.01,0.05,0.1],
eval_metric="aucpr"

{'gamma': 0, 'learning_rate': 0.05, 'max_depth': 15, 'n_estimators': 500}
grid.best_score_ = 0.94
```

```
      precision    recall  f1-score   support
0         0.94      0.98      0.96      4462
1         0.92      0.76      0.83      1237
```

```
Confusion Matrix: [[4377   85]
                   [ 302  935]]
```

```
ROC AUC Score: 0.94
```

```
best_f1: 0.83
```

```
best_thresh: 0.51
```

Stacking Classifier

```
("log", LogisticRegression(max_iter=6000, random_state=42)),
("rnd", RandomForestClassifier(max_depth=15, n_estimators=100, min_samples_split=20, min_samples_leaf=5, random_state=42)),
("xgb", XGBClassifier(n_estimators=329, learning_rate=0.05, eval_metric="aucpr", max_depth=15, random_state=42))
```

```
meta_model1=LogisticRegression(random_state=42)
```

```
      precision    recall  f1-score   support
0         0.94      0.97      0.96      4462
1         0.89      0.76      0.82      1237
```

```
Confusion Matrix: [[4349  113]
```

```

[ 291  946]]
ROC AUC Score: 0.94
best_f1: 0.83
best_thresh: 0.73

```

LightGBM

```

model=lgb.LGBMClassifier(
    n_estimators=200,
    max_depth=15,
    learning_rate=0.05,
    random_state=42
)

precision    recall  f1-score   support

0           0.93      0.98      0.96     4462
1           0.93      0.73      0.82     1237

Confusion Matrix: [[4393   69]
                   [ 331  906]]
ROC AUC Score: 0.94
best_f1: 0.82
best_thresh: 0.56

```